

# Transparent Intrusion Detection in Xen Virtual Machines using libVMI: Performance and Limitations

# Introduction to Virtualization

Virtualization :

- Containers : process-level isolation, shared with host kernel
- VMs : OS isolation, run multiple OS on the same server

Advantages :

- efficiency -> multiple services on a single server
- isolation -> improved security
- easier maintenance : pause, clone, migrate systems
- snapshot and rollback



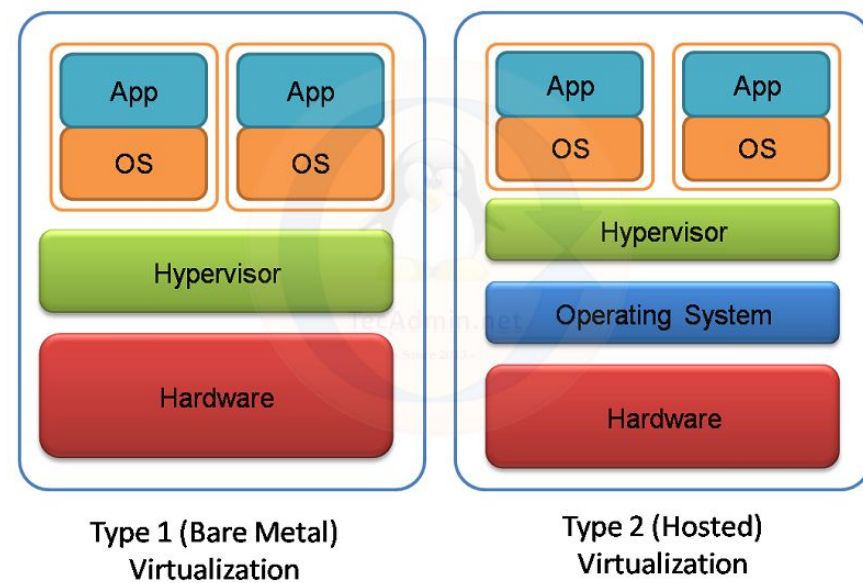
# Managing VMs

Hypervisor :

- type1: KVM, Xen, Qemu, Proxmox...
- type2 : VirtualBox, VmWare Workstation

Hypercalls :

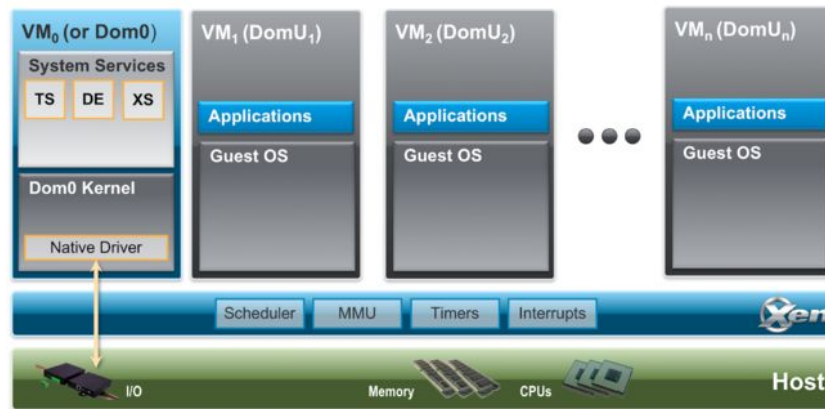
- like system calls, between VMs and hypervisor
- VM control (I/O, CPU sched, memory management)



source : <https://tecadmin.net/type-1-vs-type-2-virtualization/>

# Xen Architecture

- Xen is a type-1 (bare-metal) hypervisor:
  - Runs directly on hardware
  - Manages CPU, memory, and I/O for virtual machines
- Components:
  - **Dom0**: Privileged domain with access to hardware and control tools
  - **DomU**: Unprivileged guest VMs




source : [https://wiki.xenproject.org/wiki/Xen\\_Project\\_Software\\_Overview](https://wiki.xenproject.org/wiki/Xen_Project_Software_Overview)

# Grid5000 testbed

## Experimental Testbed for Research

- Enables fast and reproducible environment creation
- Realistic setup: literally a cluster like in production



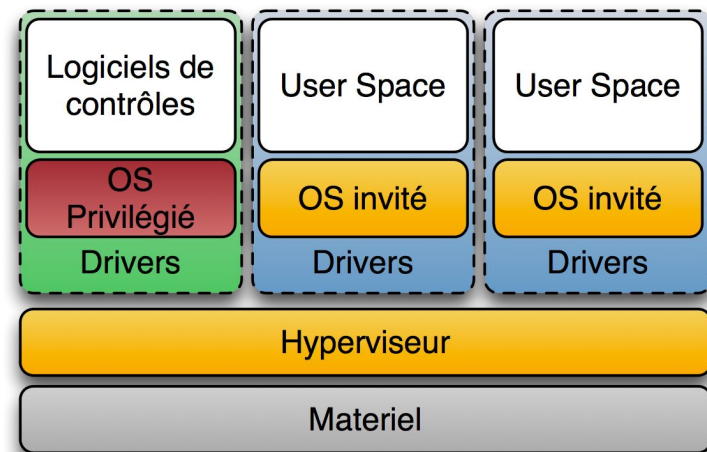
# Introduction to VMI

Architecture : Server => Hypervisor + Vms

No agents in guest => External observation

Used for IDS / malware analysis

Trade-off : security  $\leftrightarrow$  performance overhead



source : <https://fr.wikipedia.org/wiki/Hyperviseur>

# Memory introspection

libVMI - open source lib

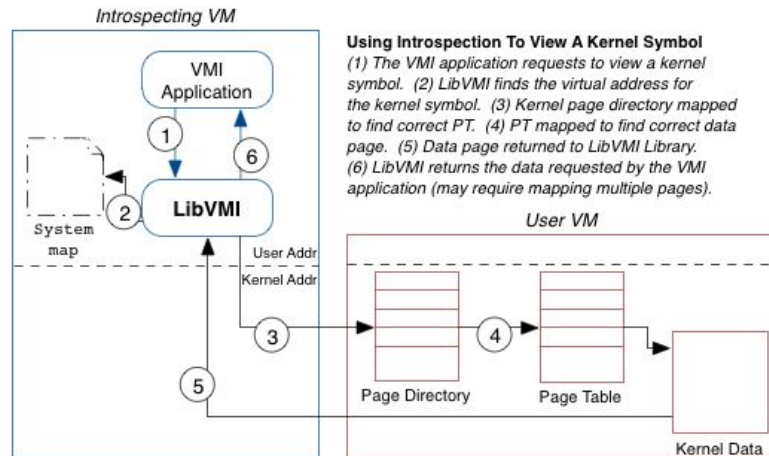
Compatible with Xen, KVM, Qemu

VMs support : Windows / Linux

32/64 bits, ARM

2 snapshot modes : live snap / halt snap

Requires : /boot/System.map (from VM), and the offsets (kernel module => [findoffets.so](https://findoffets.so))



source : <https://libvmi.com/docs/gcode-intro.html>



# Benchmark tools & methodology

Tools :

- Phoronix Test Suite
- perf
- htop & xl top
- xentrace & xenalyze => custom parser + sync script

For every bench of Phoronix :

- with libVMI
- without libVMI





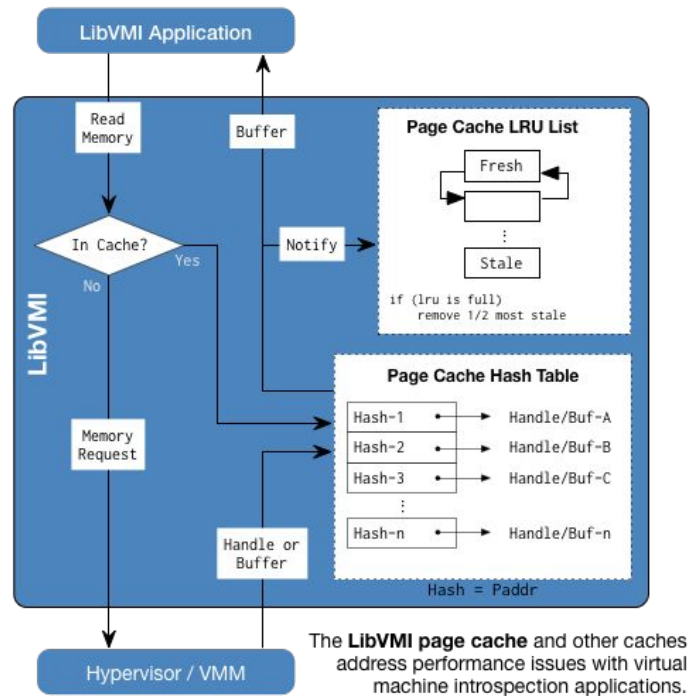
# Init Setup & Cache

Init setup :

- vmi-process-list
- halt mode, external loop

Cache :

- Least Recently Used
- Purpose: Speeds up memory access
- Method: Caches to avoid repeated physical reads
- Improves memory analysis efficiency
- Validity: Only valid when VM is running or paused



Children	Self	Command	Shared Object	Symbol
+ 87.23%	0.00%	vmi-process-lis	libc.so.6	[.] 0x00007fa9912fe24a
+ 87.23%	0.00%	vmi-process-lis	vmi-process-list	[.] main
+ 86.92%	0.00%	vmi-process-lis	libvmi.so.0.0.15	[.] linux_symbol_to_address
+ 86.92%	0.00%	vmi-process-lis	libvmi.so.0.0.15	[.] linux_system_map_symbol_to_address
+ 67.88%	0.00%	vmi-process-lis	libvmi.so.0.0.15	[.] vmi_init_complete

# Optimization

1. init overhead removed
  - overhead found with perf (dom0)
  - 0,30s in vmi\_init\_complete (clock time)
2. disabled VM pausing
3. improved robustness
  - max 5 iterations
4. enable continuous introspection
5. added page cache flushing
  - for reboot / VM delition
  - cache becomes obsolete, causing errors or inconsistencies




# Performance impact in DomU


Benchmarks used :

- smallpt
- compress-gzip
- c-ray (4k)
- build-linux-kernel (defconfig)
- apache (200), pts/stream


LibVMI overhead depending on version



# Most Optimized detailed results



# Most accurate detailed results



# Profiling limitations

## DomU profiling :

- limited (perf) => paused during introspection with “initial version”
- 2.40% overhead (with init version) => recovery cost after pause

## Dom0 profiling :

- htop : no accurate data

## Hypervisor profiling :

- xl top
- libVMI : dom0 39% usage, domU 740% CPU usage
- no libVMI : dom0 1.8%, domU 800%



# Xentrace & Xenalyze

Tracing :

- xentrace -> raw binary file
- xenalyze : readable trace “summary”
- custom python parser : extract and generate stats

Methods :

- VM idle
- libVMI only
- Bench only
- libVMI + bench



# Parser output

## Findings :

- compress-gzip bench : 30% overhead
- libVMI with cache flush
- domU : hypercalls +15%
- dom0 : increased significantly  
(iret, stack\_switch, mmu\_update, vcpu\_op)
- stable VM CPU time



```
Statistics from traces/trace_nolibvml.txt -----
Total tracing time : 8.72 seconds

Domain 0: 64 vCPUs
Running      mean=0.00s, median=0.00s, min=0.00s, max=0.07s
Runnable     mean=0.00s, median=0.00s, min=0.00s, max=0.00s
Preempt      (no data)
Blocked      mean=5.96s, median=6.37s, min=3.49s, max=8.71s
Wake         mean=0.00s, median=0.00s, min=0.00s, max=0.00s
Hypercalls   mean=208, median=56, min=24, max=4479
PTWR         mean=4, median=4, min=4, max=4
Privop Emu   mean=135, median=6, min=4, max=6819
Hypercall types:
- iret       : 4834
- vcpu_op    : 3177
- set_segment_base : 1830
- stack_switch : 1829
- sched_op   : 1454
- evtchn_op  : 142
- grant_table_op : 29
- mmuext_op  : 24
- xen_version : 6
- mmu_update : 5
- physdev_op : 2
- sysctl     : 1

Domain 3: 8 vCPUs
Running      mean=0.00s, median=0.00s, min=0.00s, max=0.01s
Runnable     mean=0.00s, median=0.00s, min=0.00s, max=0.00s
Preempt      (no data)
Blocked      mean=6.57s, median=7.25s, min=3.94s, max=8.66s
Wake         mean=0.00s, median=0.00s, min=0.00s, max=0.00s
Hypercalls   mean=149, median=69, min=38, max=618
PTWR         (no data)
Privop Emu   mean=151, median=194, min=26, max=232
Hypercall types:
- iret       : 435
- vcpu_op    : 363
- sched_op   : 209
- stack_switch : 150
- evtchn_op  : 25
- set_segment_base : 8
- mmuext_op  : 3

Domain 32767: 47 vCPUs
Running      mean=6.88s, median=7.79s, min=4.49s, max=8.71s
Runnable     mean=0.00s, median=0.00s, min=0.00s, max=0.07s
Preempt      mean=0.00s, median=0.00s, min=0.00s, max=0.07s
Blocked      (no data)
Wake         (no data)
Hypercalls   (no data)
PTWR         (no data)
Privop Emu   (no data)
```



# Conclusion

- libVMI enables monitoring, but at a cost
- removing VM pause reduces performance overhead
- cache flushing ensures accuracy, but reintroduces overhead
- traditional tools miss indirect effects of introspection : hypervisor-level tools required

Future work :

- still a “black box”
- hard to get useful data from Xen
- adapt introspection frequency
- adapt cache flush frequency

